

Content



0. Introduction

1. Regression

1.1 Multivariate Linear Regression (curve fitting)

1.2 Regularization (Lagrange multiplier)

1.3 Logistic Regression (Fermi-Dirac distribution)

1.4 Support Vector Machine (high-school geometry)

2. Dimensionality Reduction/feature extraction

2.1 Principal Component Analysis (order parameters)

2.2 Recommender Systems

2.3 Clustering (phase transition)

Content



3. Neural Networks

3.1 Biological neural networks

3.2 Mathematical representation

3.3 Feed-forward neural networks

3.5 Different Learning algorithm

3.6 Deep learning and CNN

Improving the way neural networks learn

- A better choice of cost function — the cross-entropy
- Regularization methods — L1, L2, dropout, Jackknife resampling
- A better initialisation of the weights
- Good choice of hyper-parameters

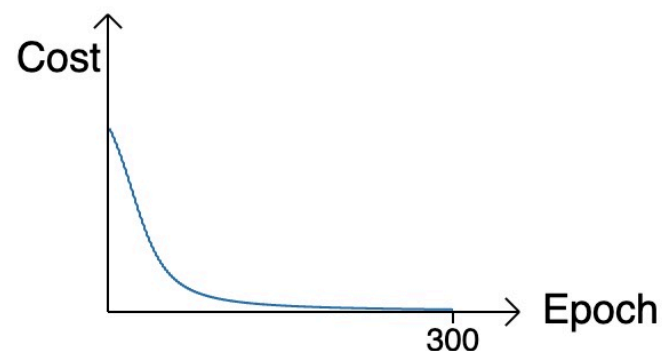
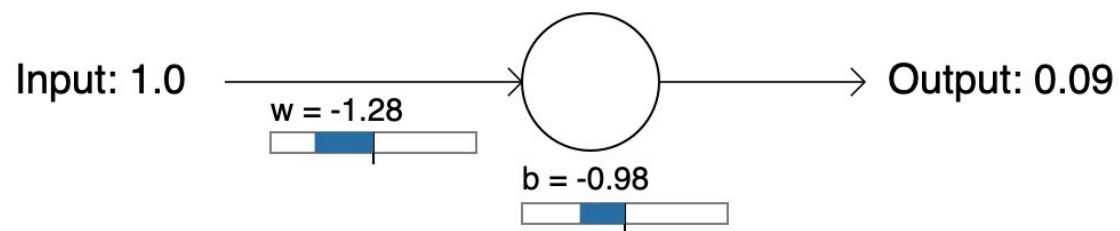
We often learn fastest when we're badly wrong about something.
But our artificial neuron has a lot of difficulty learning when it's badly wrong - far more difficulty than when it's just a little wrong.

Starting $w = 0.6$, $b = 0.9$

input=1, expected output 0, initial output 0.82

quadratic cost

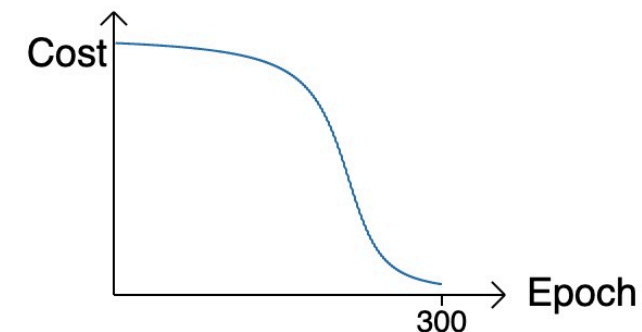
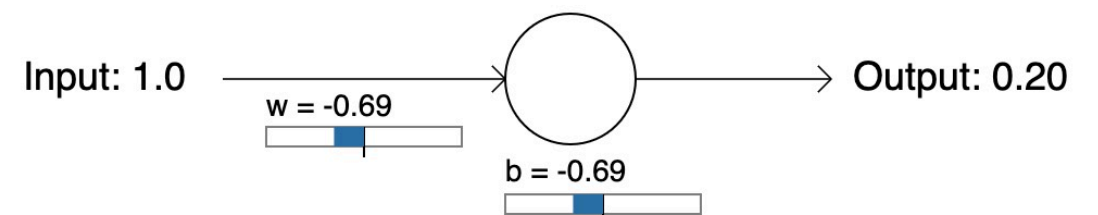
$$J = \frac{1}{2}(y - a)^2$$



Run

Starting $w = 2$, $b = 2$

input=1, expected output 0, initial output 0.98



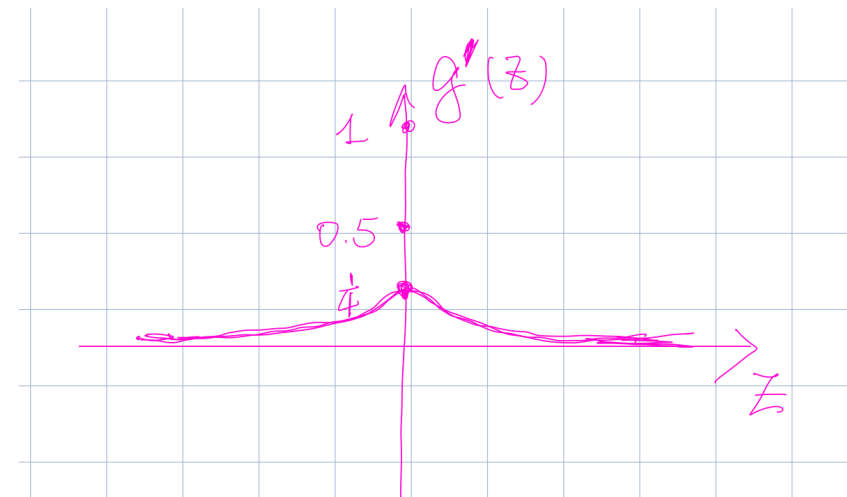
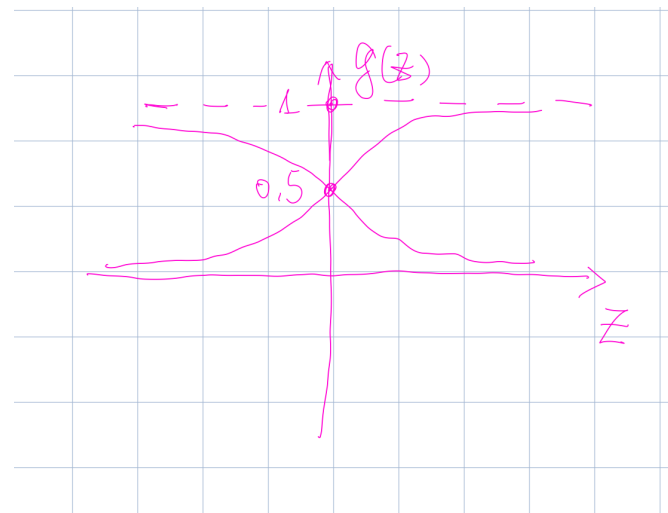
Run

Learning slow down at the saturation

$$\frac{\partial J}{\partial w} = (a - y)g'(z)x = ag'(z)$$

$$\frac{\partial J}{\partial b} = (a - y)g'(z) = ag'(z)$$

when $x = 1$ and $y = 0$

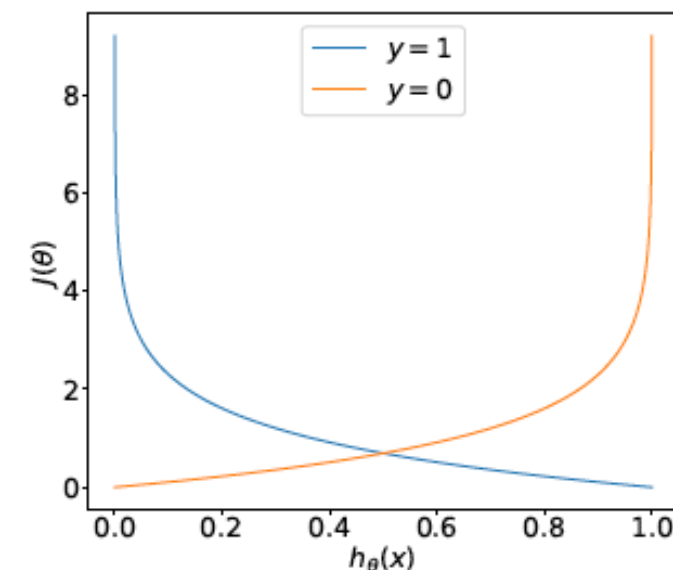


$$\frac{\partial}{\partial \Theta_{ij}^l} J(\Theta) = a_j^l \delta_i^{l+1} \quad \delta^3 = (\Theta^3)^T \delta^4 \cdot * g'(z^3) \quad \text{Learning slowdown when output neuron saturates}$$

Cross-entropy cost function

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m [y_i \ln(a_i^L) + (1 - y_i) \ln(1 - a_i^L)]$$

- the cross-entropy is non-negative
- if the neuron's actual output is close to the desired output for all training inputs the cross-entropy will be close to zero
- avoid the learning slow down



$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m [y_i \ln(a_i^L) + (1 - y_i) \ln(1 - a_i^L)]$$

$$\frac{\partial J}{\partial \theta_i} = \frac{1}{m} \sum_{i=1}^m x_i (g(z_i) - y_i)$$

Error in output control the learning

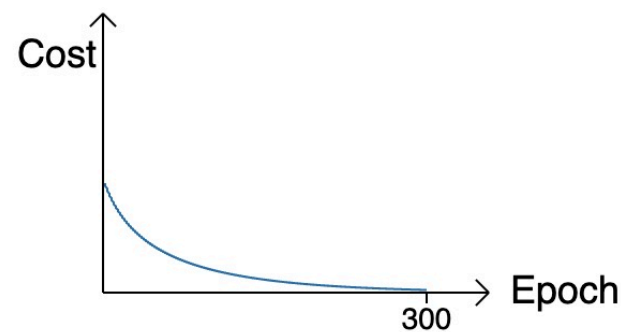
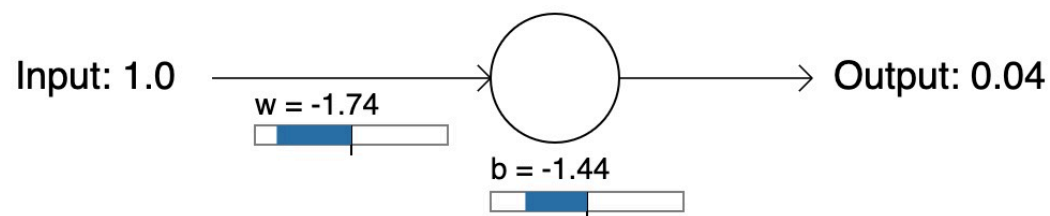
$$\frac{\partial J}{\partial b} = \frac{1}{m} \sum_{i=1}^m (g(z_i) - y_i)$$

(remember gradient in logistic regression ?)

Starting $w = 0.6, \quad b = 0.9$

cross-entropy cost

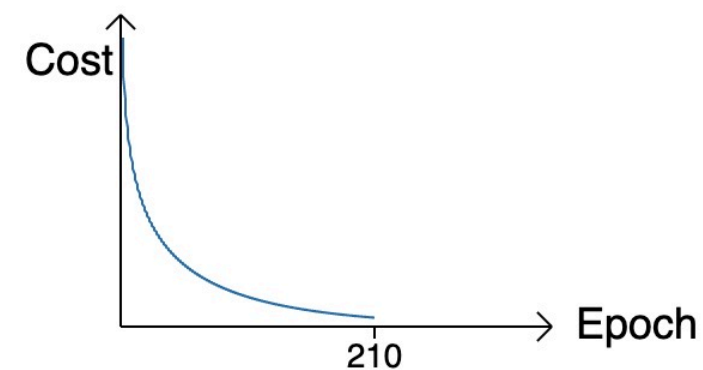
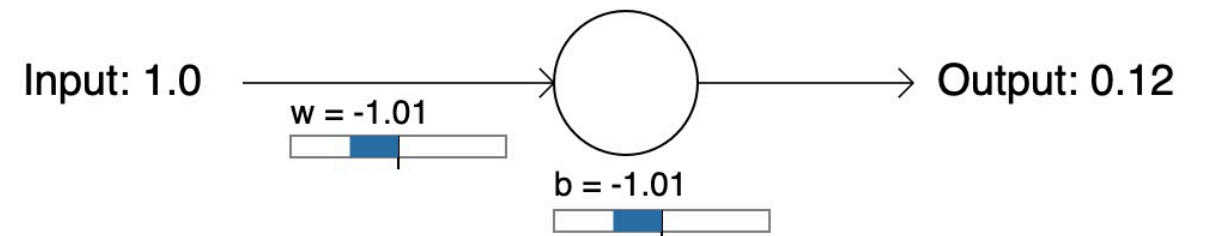
input=1, expected output 0, initial output 0.82



Run

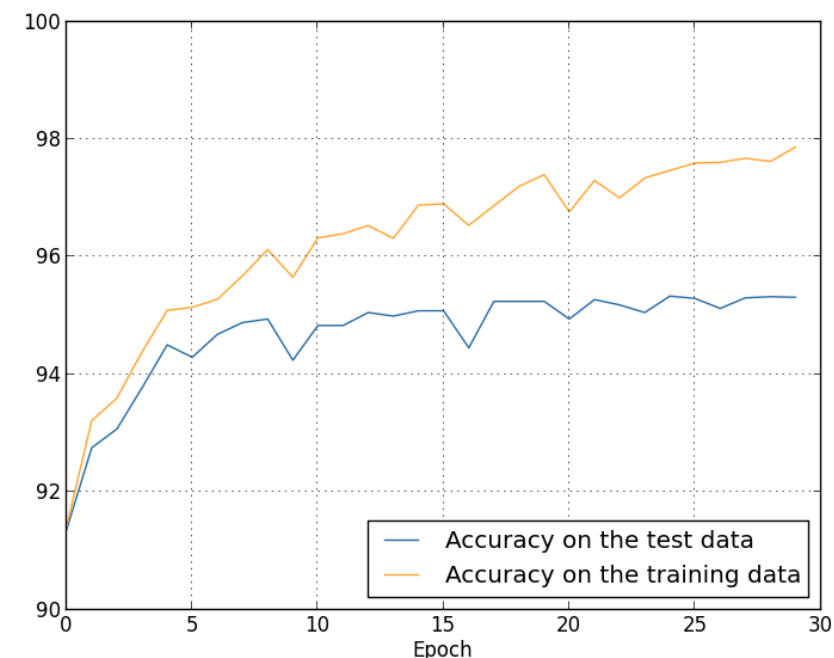
Starting $w = 2, \quad b = 2$

input=1, expected output 0, initial output 0.98



- when we use the quadratic cost learning is *slower* when the neuron is wrong than it is later on, as the neuron gets closer to the correct output
- while with the cross-entropy learning is faster when the neuron is wrong.
- from an information perspective, the cross-entropy measures how "surprised" we are, on average, when we learn the true value for y . We get low surprise if the output is what we expect, and high surprise if the output is unexpected

50000 training
10000 test
30 hidden neutrons, mini_batch_size=10 $\eta = 0.5$



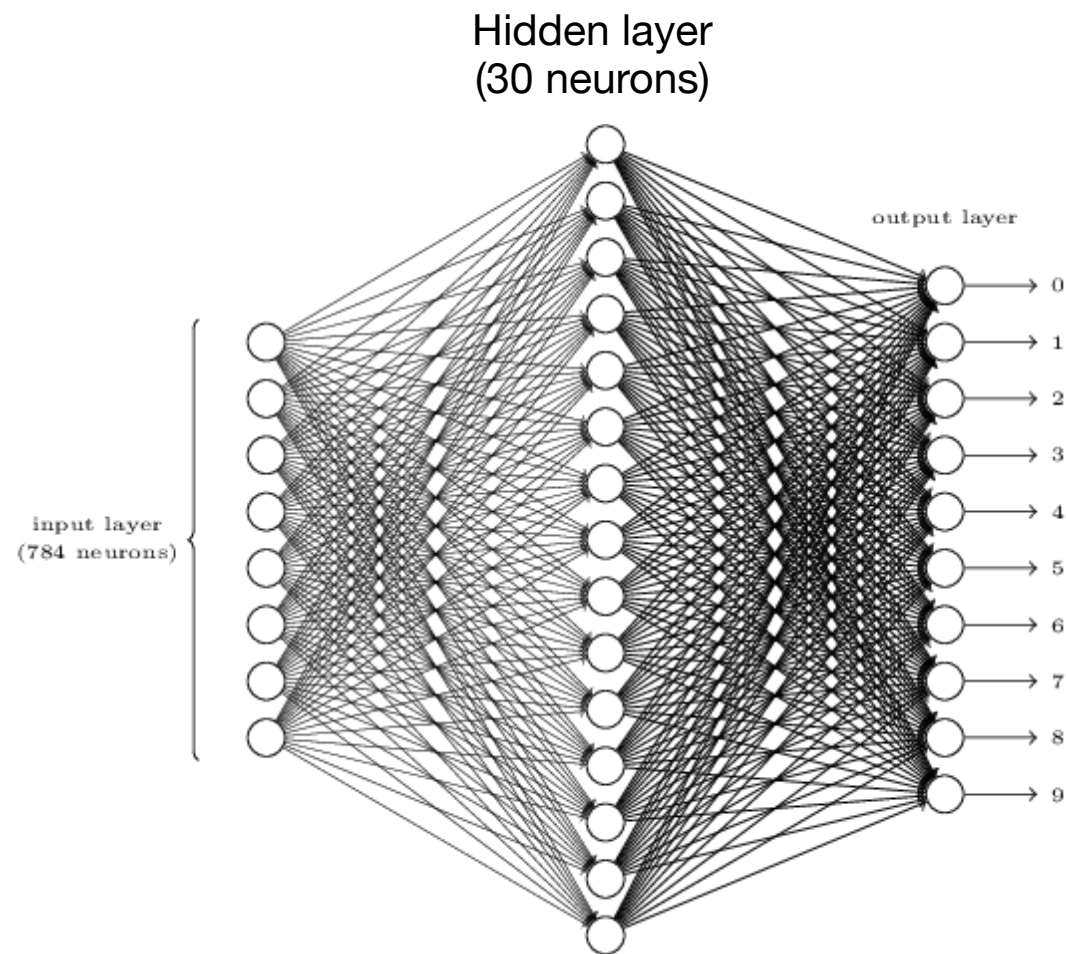
Regularization

$$J(\Theta) = -\frac{1}{m} \sum_{i=1}^m [y_i \ln(a_i^L) + (1 - y_i) \ln(1 - a_i^L)] + \frac{\lambda}{2m} \sum_{\Theta} \Theta^2$$

$$\Theta = (1 - \frac{\eta\lambda}{m})\Theta - \frac{\eta}{m} \sum_i \frac{\partial J}{\partial \Theta}$$

Hyperparameters: learning rate, mini-batch size

Requires thorough optimisation process



Fermi: "I remember my friend Johnny von Neumann used to say, with four parameters I can fit an elephant, and with five I can make him wiggle his trunk."

Training_data

validation_data

test_data

$$\{\Theta_{ij}^{(l)}\}$$

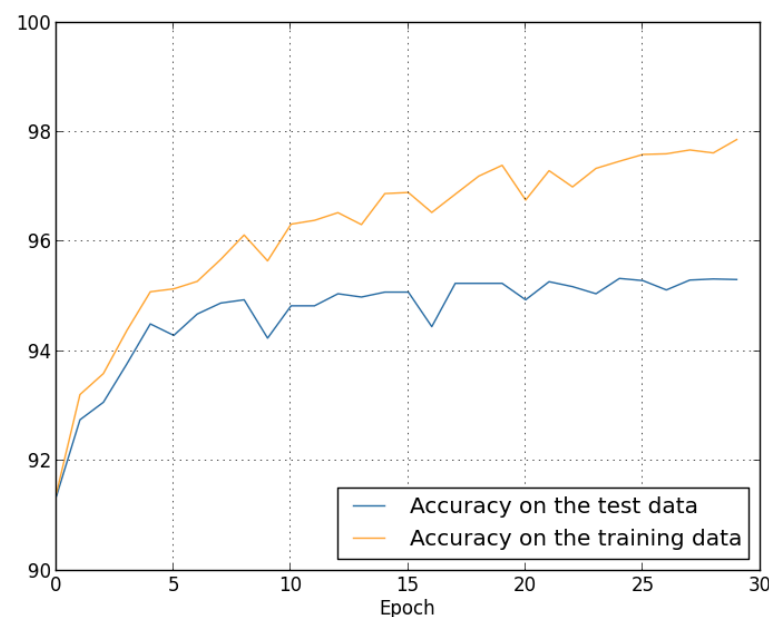
Hyperparameters

50,000 images
30 epochs, mini-batch size 10

10,000 images

10,000 images

$$785 \times 30 + 10 \times 31 = 23,860$$



Without regularisation



With regularisation

the validation data as a type of training data that helps us learn good hyper-parameters.

50000 training
10000 test
30 hidden neutrons, mini_batch_size=10
 $\eta = 0.5, \quad \lambda = 5$

- choice a better activation — softmax

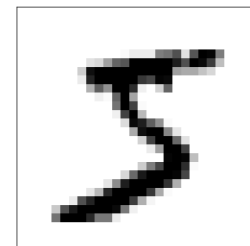
Softmax activation — partition function $a_j^L = \frac{e^{z_j^L}}{\sum_k e^{z_k^L}}$

- Regularization methods — L1, L2, dropout,

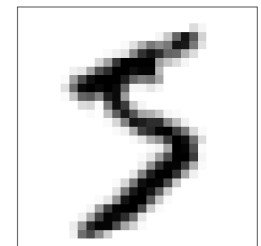
$$J(\Theta) = J(\Theta) + \frac{\lambda}{m} \sum_{\Theta} |\Theta|$$

$$J(\Theta) = J(\Theta) + \frac{\lambda}{m} \sum_{\Theta} \Theta^2$$

- Jackknife resampling



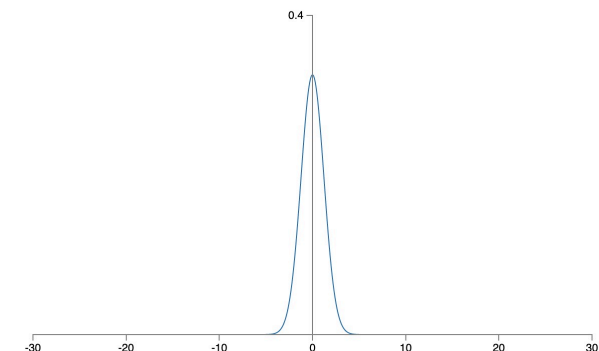
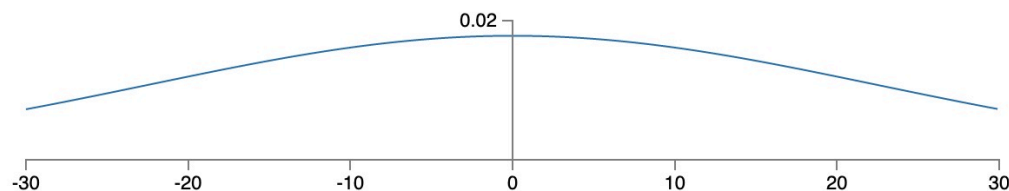
rotate/distort



- A better initialisation of the weights — do not saturate the neurons

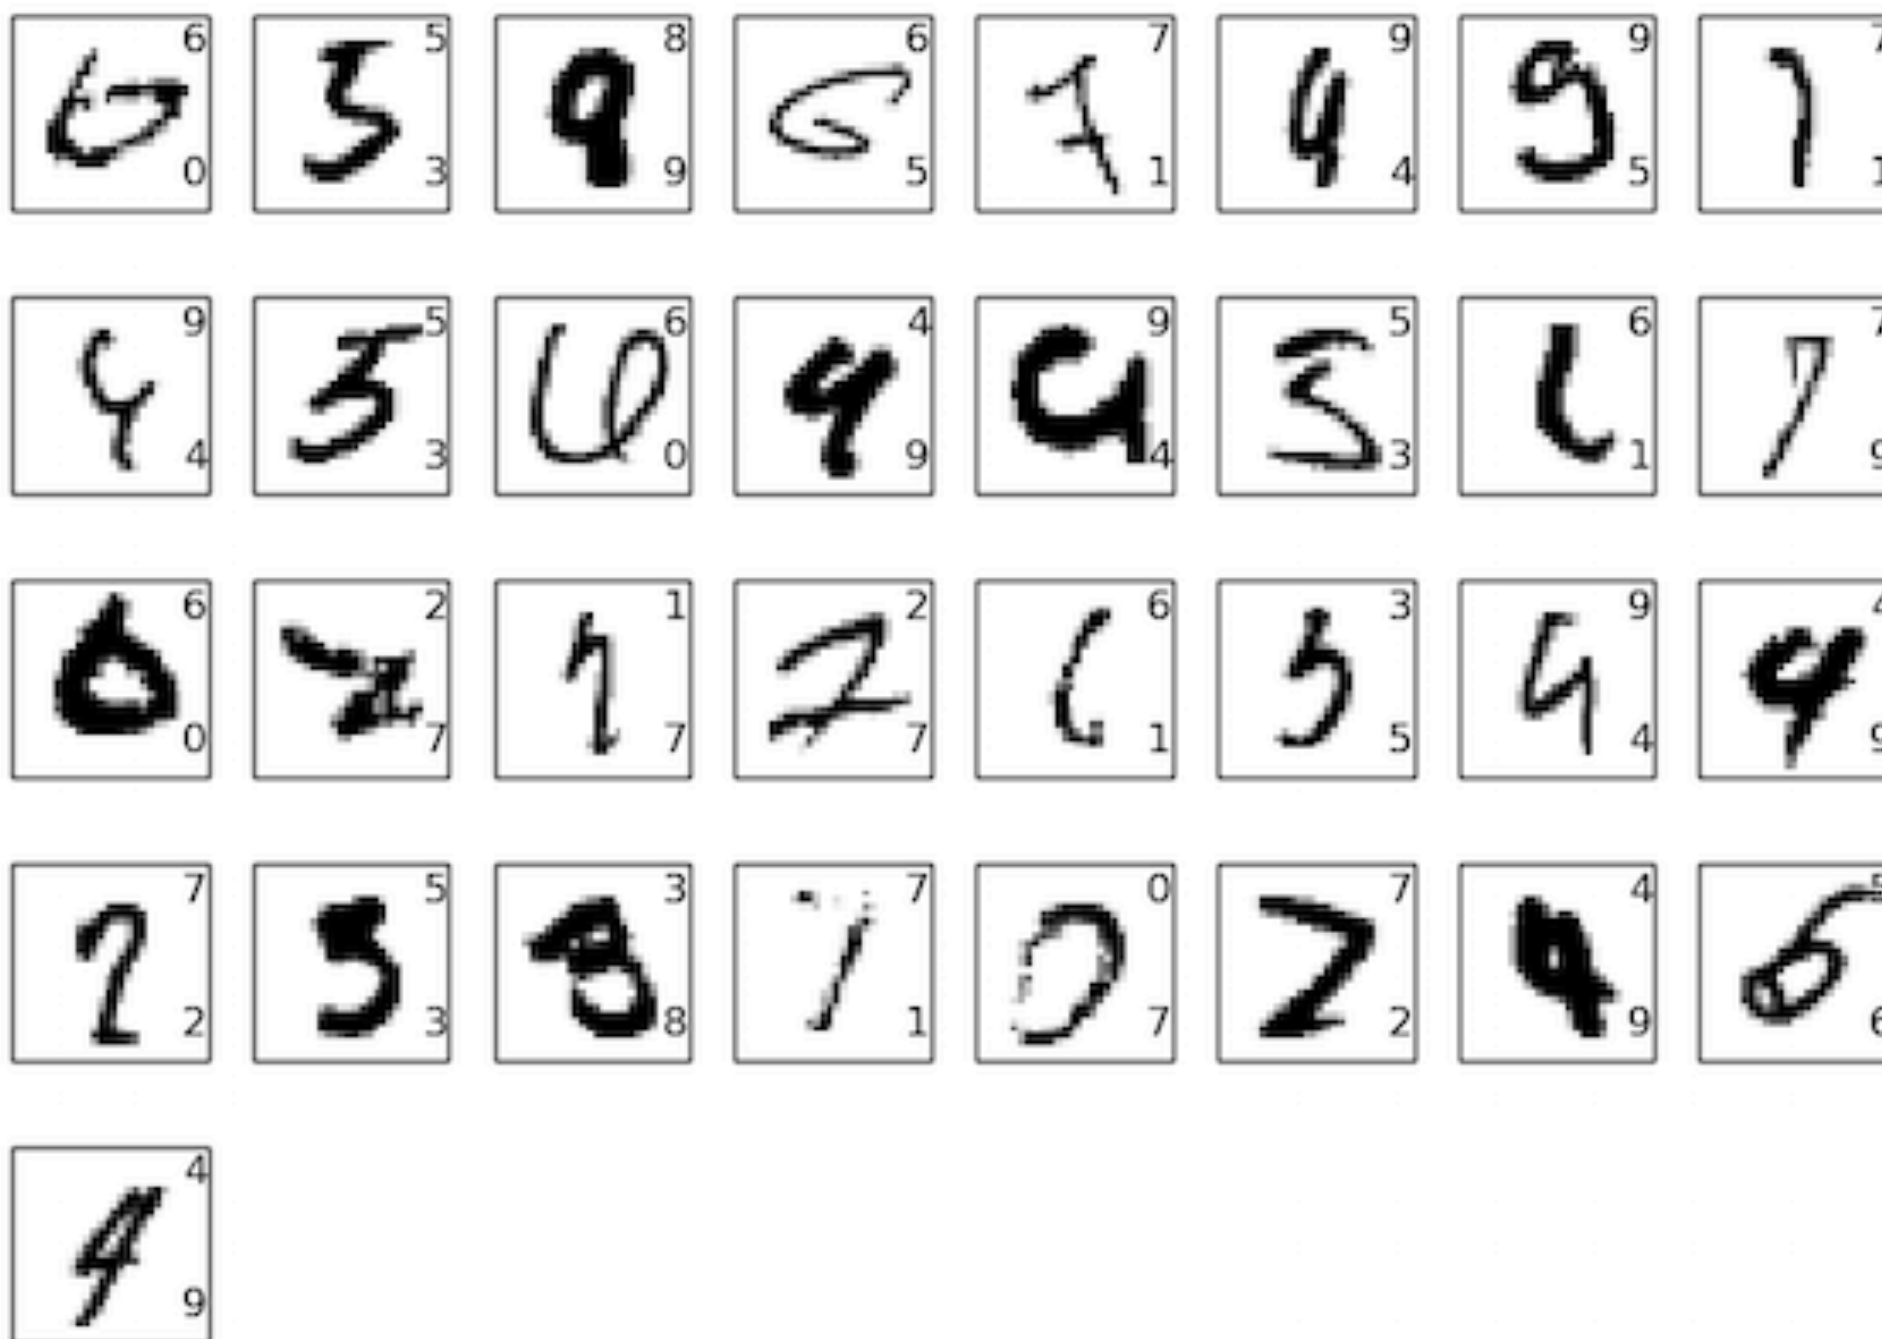
$$z = \sum_i^{n_{in}} \Theta_i x_i + b \quad \Theta \in N(0,1) \quad z \gg 1 \quad \textbf{or} \quad z \ll -1$$

$$\Theta \in N(0, 1/\sqrt{n_{in}})$$



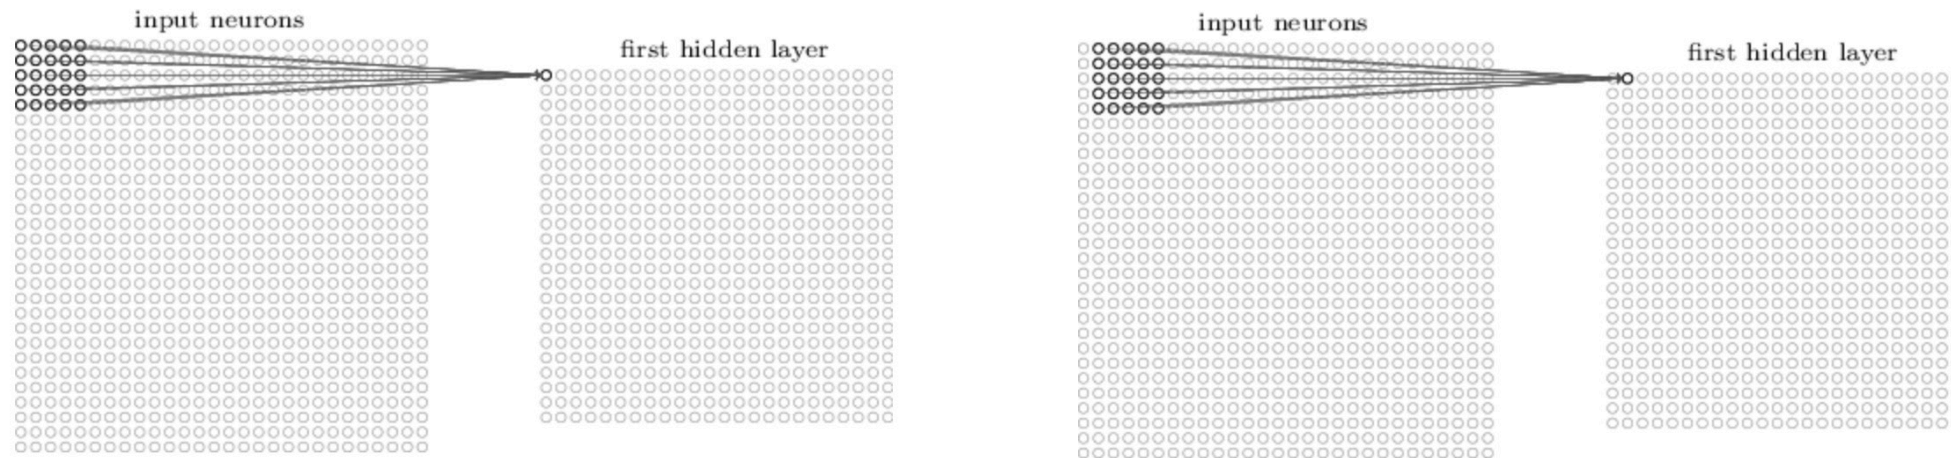
Deep learning - Convolutional neural network (CNN)

With CNN, 10,000 MNIST test images, one can classify 9,967 correctly. Accuracy 99.67% !!
Here's a peek at the 33 images which are misclassified. These are tough even for humans.



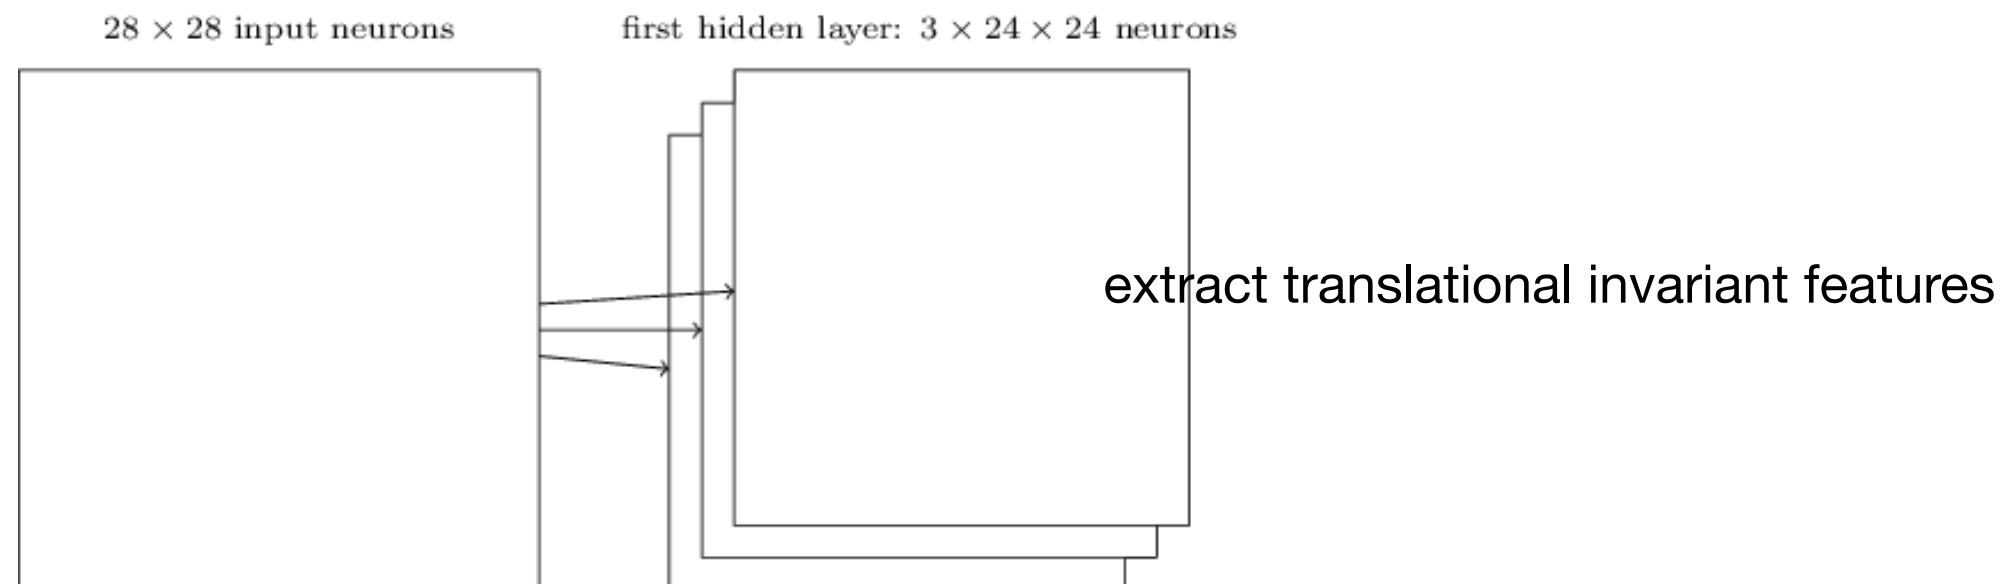
Deep learning - Convolutional neural network (CNN)

5×5 **local receptive fields** the same weights and bias $g(b + \sum_{l=0}^4 \sum_{m=0}^4 \Theta_{l,m} a_{j+l,k+m})$ $j = 1, 2, \dots, 24$
 $k = 1, 2, \dots, 24$



Convolutional layer

More than one **feature map**, more than one set of shared weights+bias, more than one **filter**



LeNet-5, 6 feature maps, 5x5 local receptive field

Y. Lecun, L. Bottou, Y. Bengio and P. Haffner, "[Gradient-based learning applied to document recognition](#)"

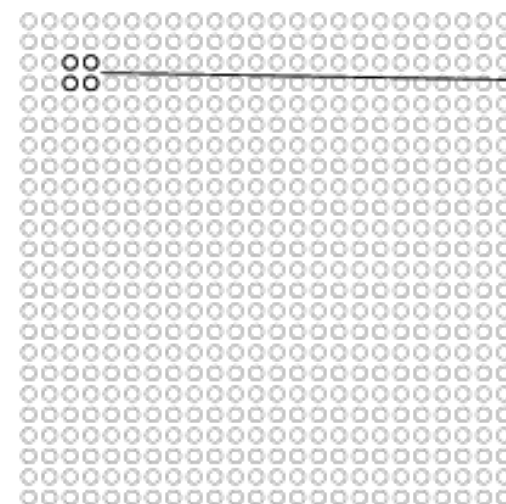
Proceedings of the IEEE, vol. 86, no. 11, pp. 2278-2324 (1998)

Reduced number of parameters:
Each feature map 5×5 (weights) + 1 (bias) = 26 parameter
20 feature maps $20 \times 26 = 520$ parameters in total

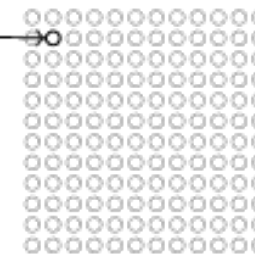
$785 \times 30 = 23550$ fully-connected layer

Pooling layers max-pooling, take the maximum activation in the 2×2 input region

hidden neurons (output from feature map)



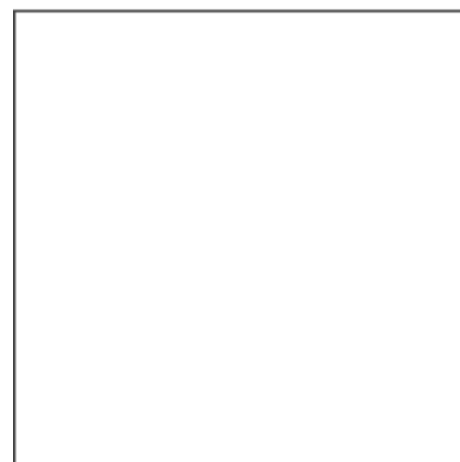
max-pooling units



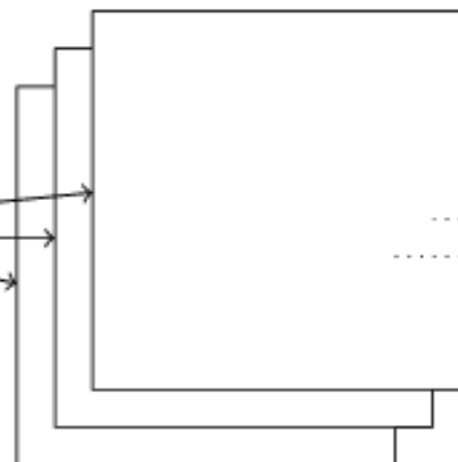
12×12

24×24

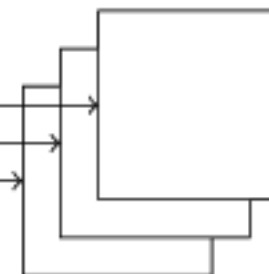
28×28 input neurons

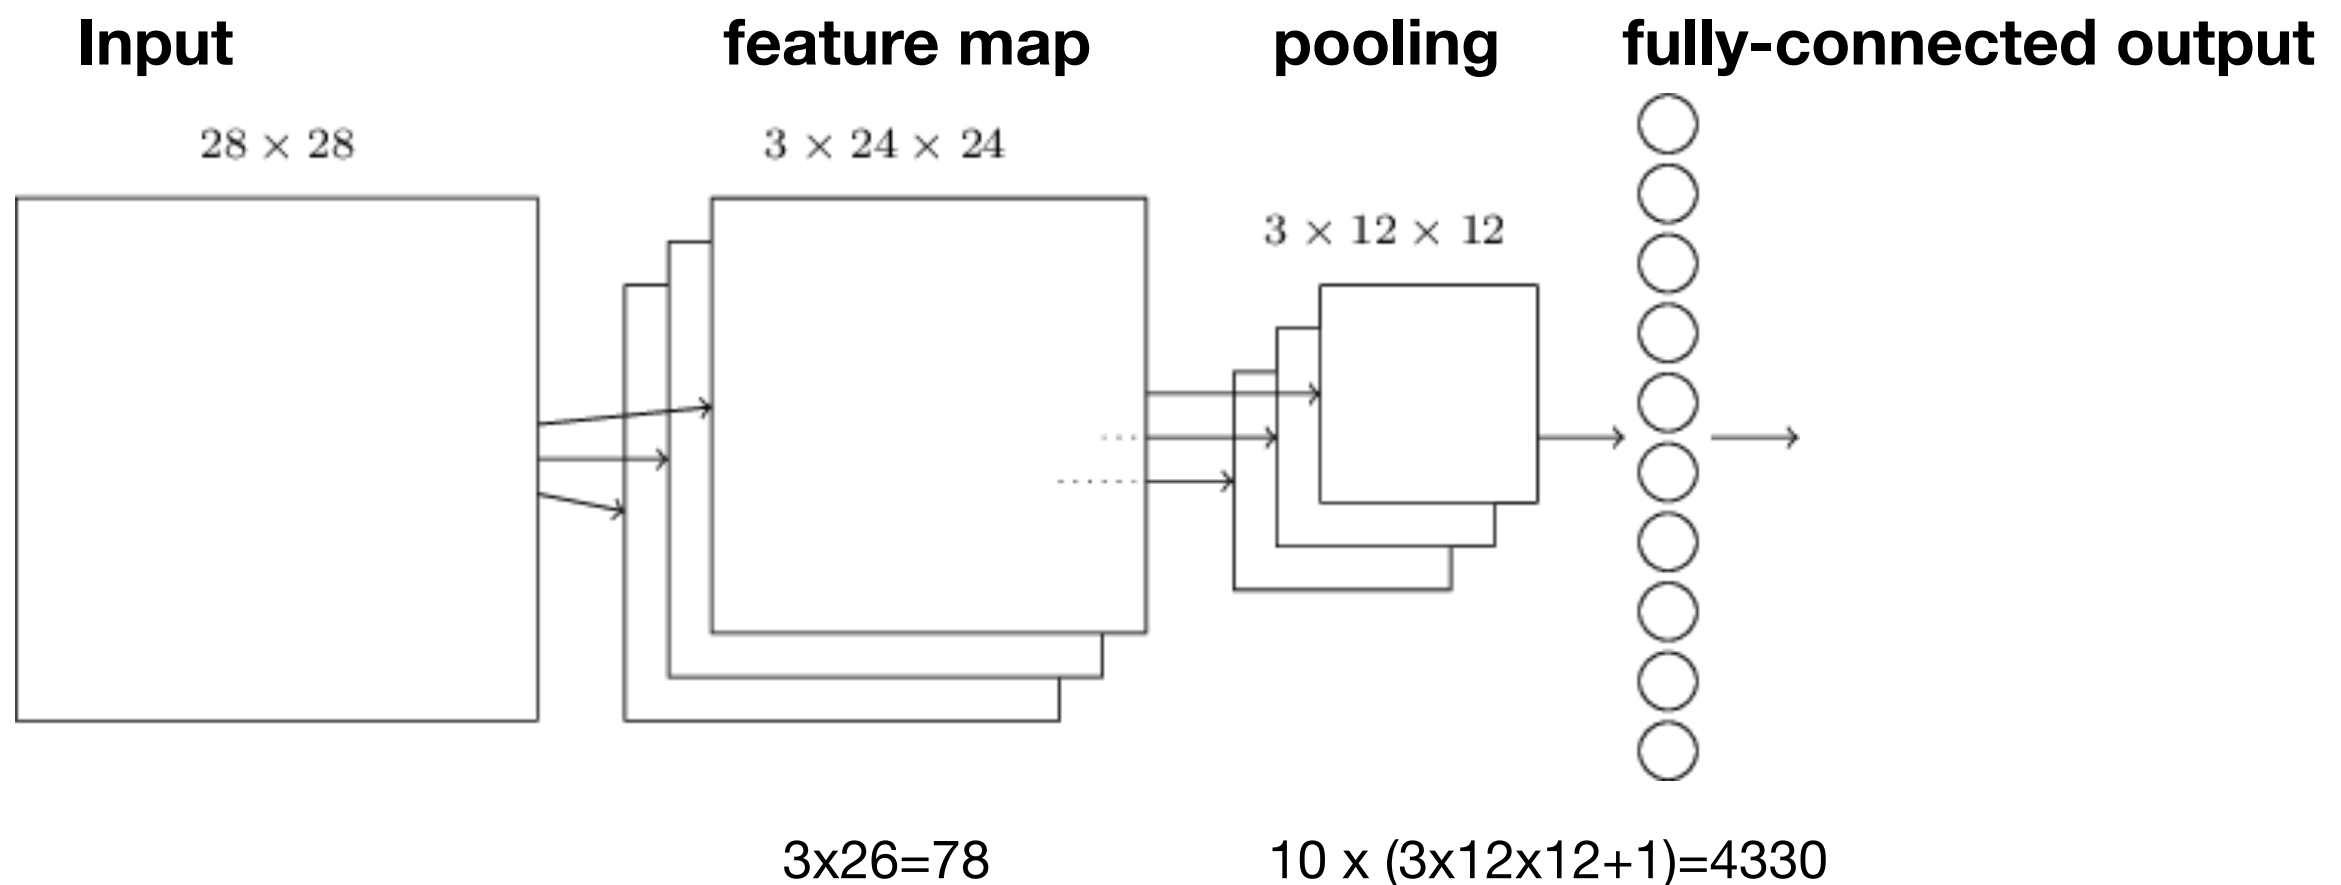


$3 \times 24 \times 24$ neurons



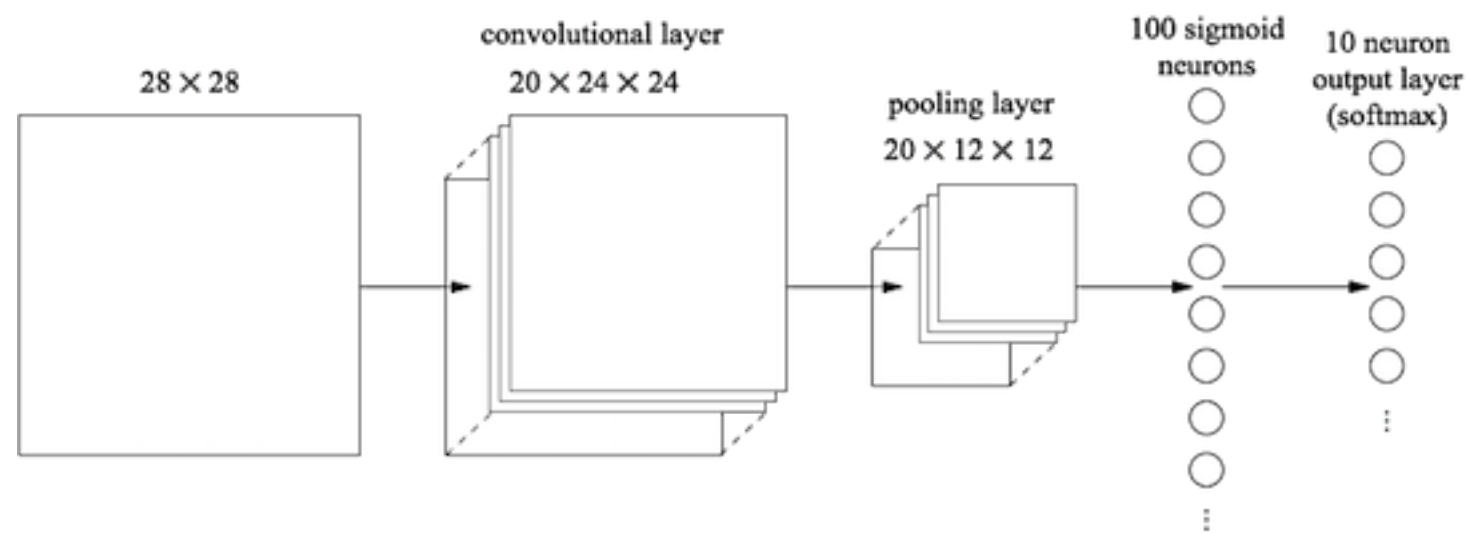
$3 \times 12 \times 12$ neurons





Backpropagation for pooling and convolution layers

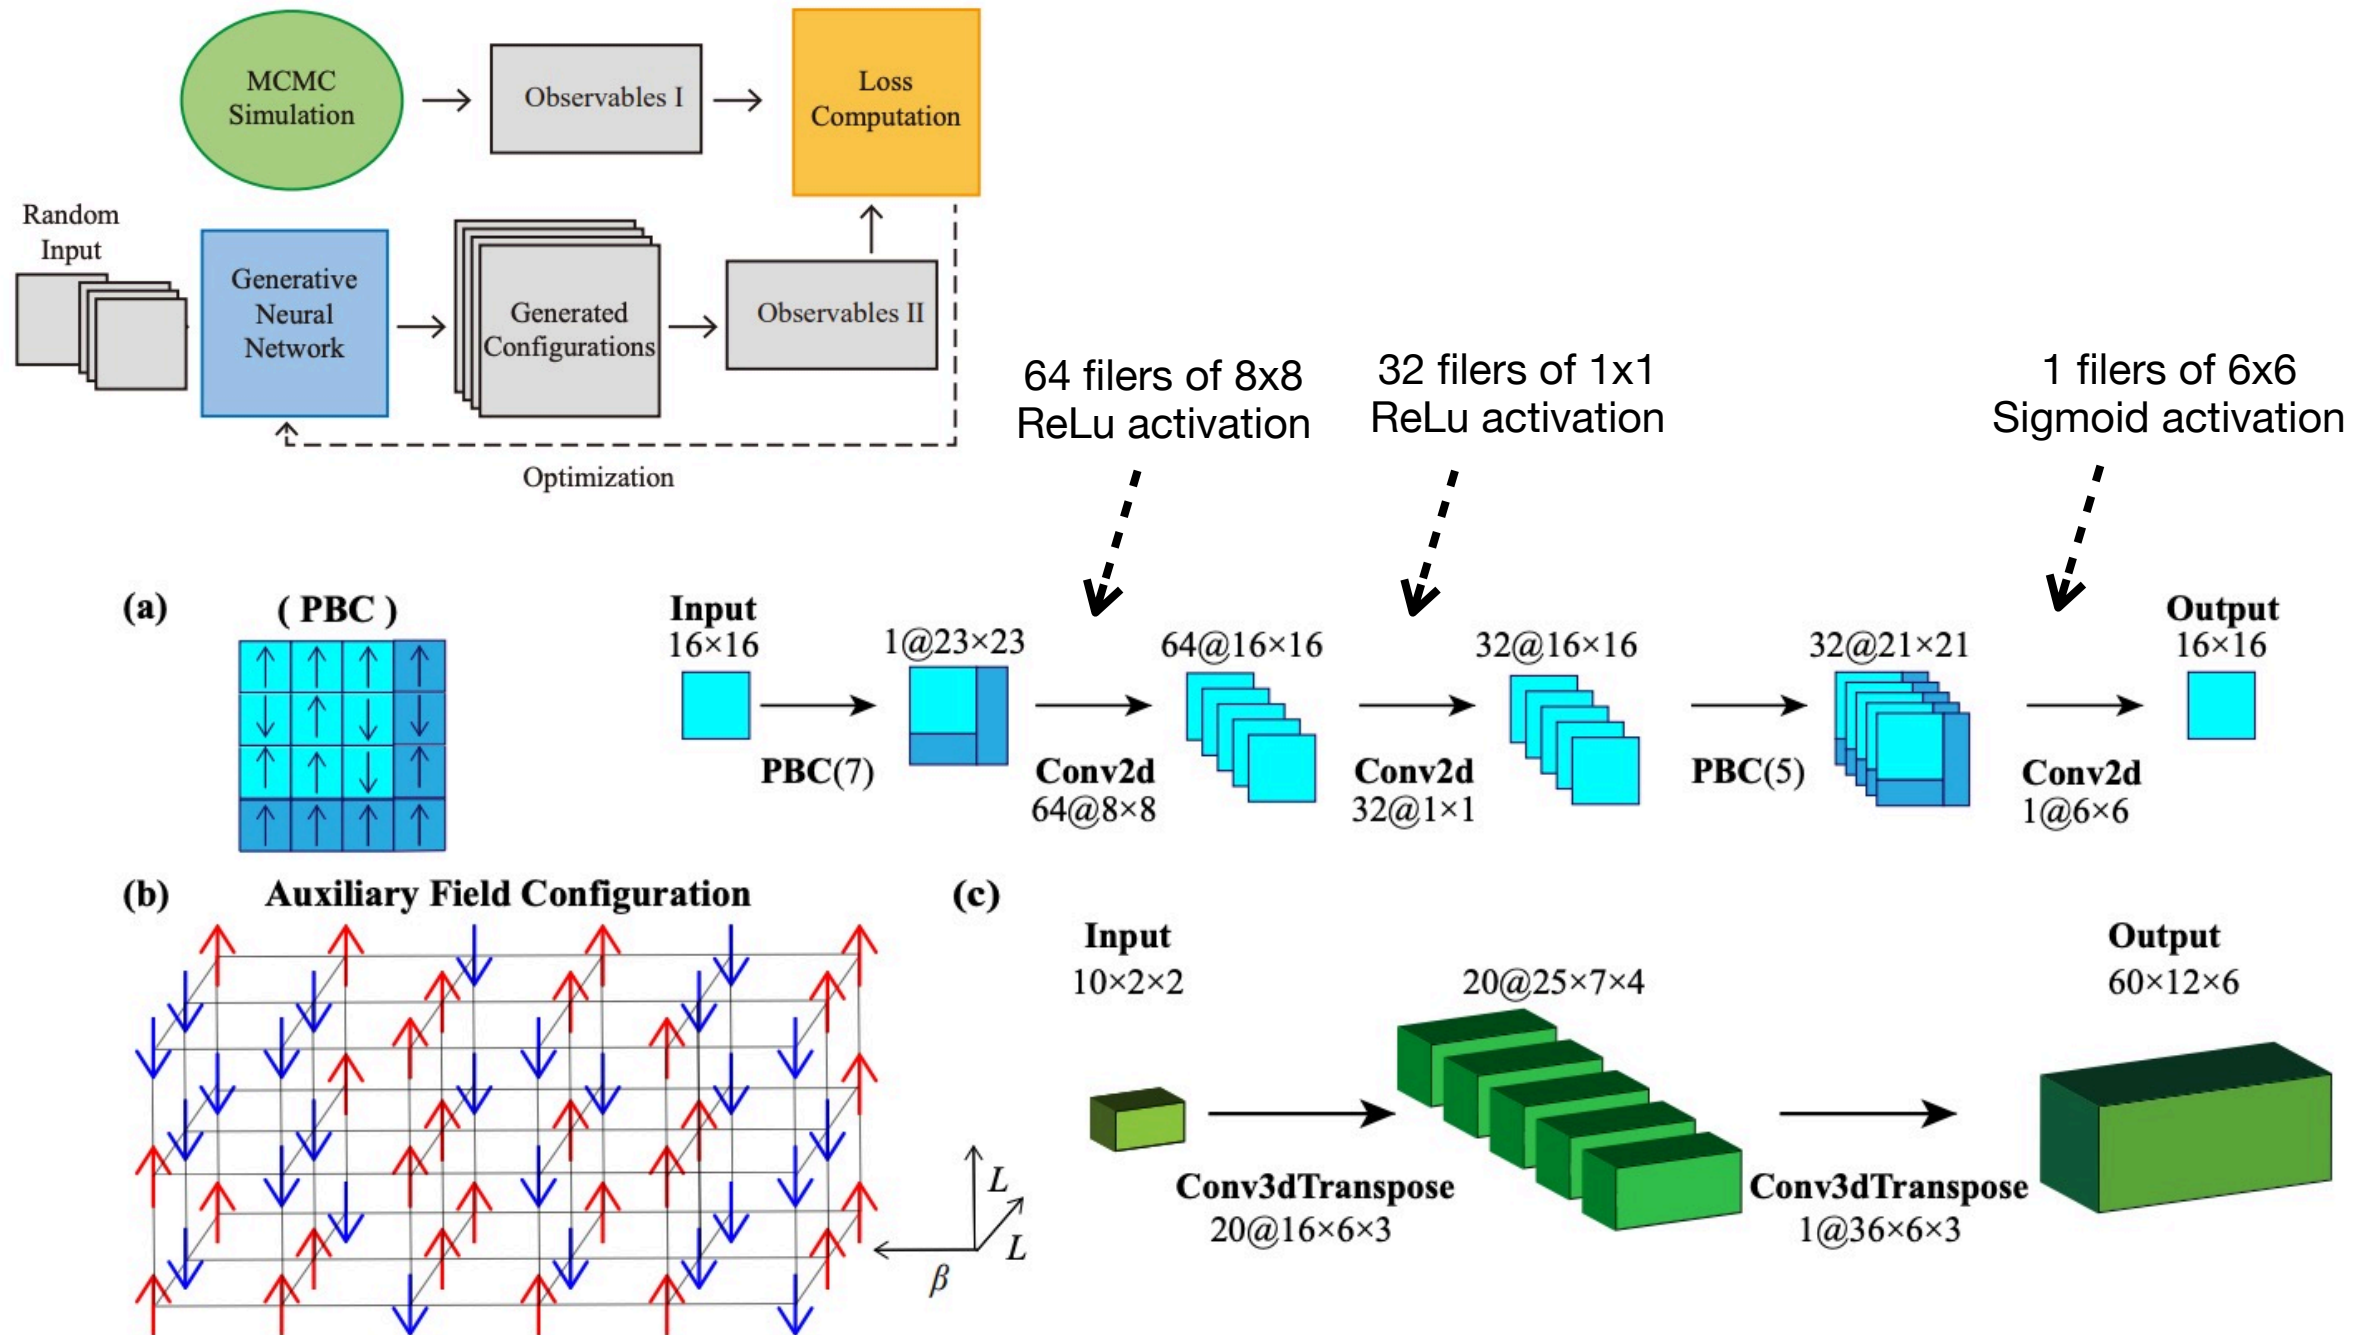
This architecture



gives 98.78 % accuracy.

Network-Initialized Monte Carlo Based on Generative Neural Networks

Hongyu Lu, Chuhao Li, Bin-Bin Chen, ..., ZYM

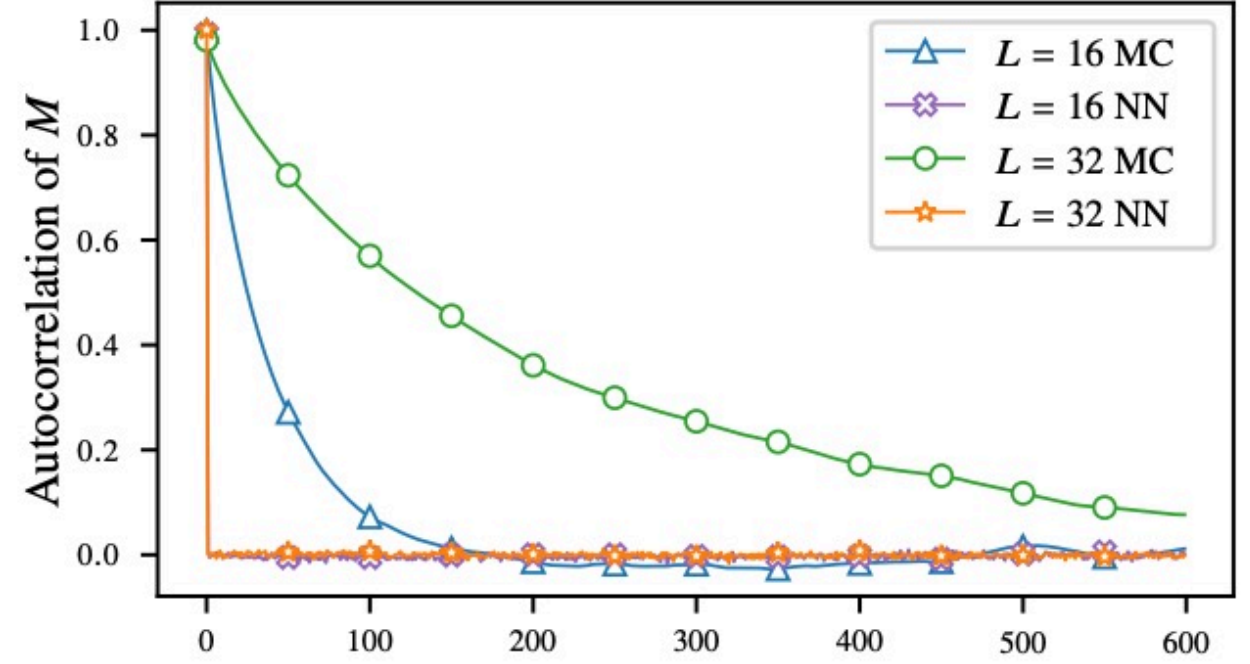


In order to optimize such neural networks, for both classical and quantum models, we prepare 1000 sets of observables measured from MC simulations and take 1000 input random configurations which will be processed into generated configurations. The comparison in loss function is randomly distributed without any grouping. For the choice of observables in the defined loss functions, we simply pick some from the ones we usually focus on. In the Ising case, we take the batch size to be 5 and epoch number to be up to 150 as the computation is quite easy. While in the Hubbard case, we run at most 15 epochs with a batch size of 3. We optimize the network parameters using Adam[63] with conventional learning rate 10^{-3} , $\beta_1 = 0.9$, and $\beta_2 = 0.999$.

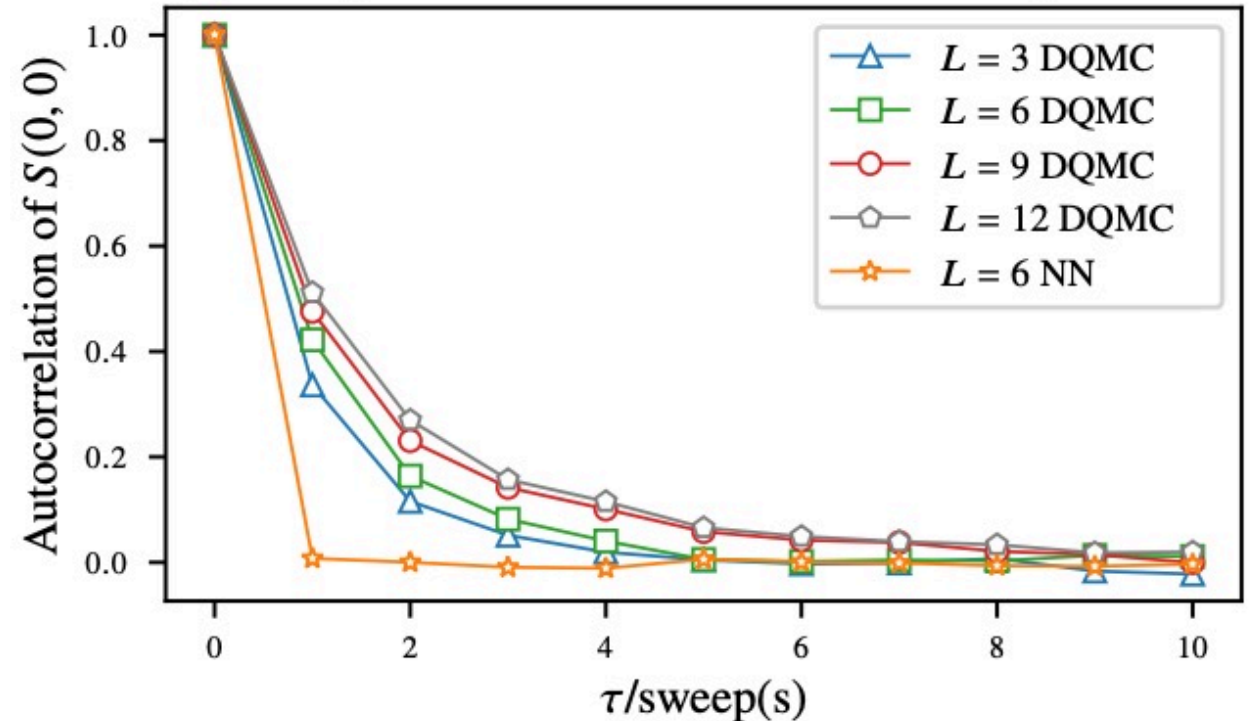
$$\begin{aligned} \text{Loss}(\mathbf{G}_l; M_l, E_l) = & w_1 \sum_l [|\mathbf{M}(\mathbf{G}_l)| - |M_l|]^2 \\ & + w_2 \sum_l [M^2(\mathbf{G}_l) - M_l^2]^2 + w_3 \sum_l [E(\mathbf{G}_l) - E_l]^2 \end{aligned} \quad (2)$$

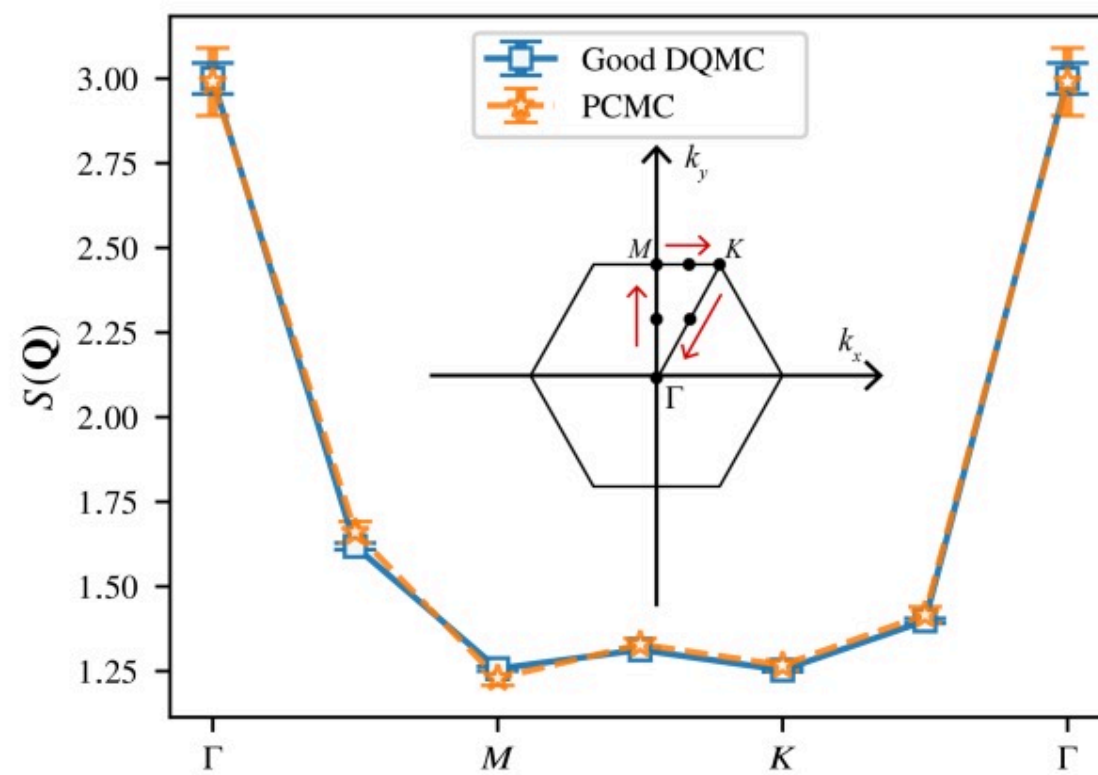
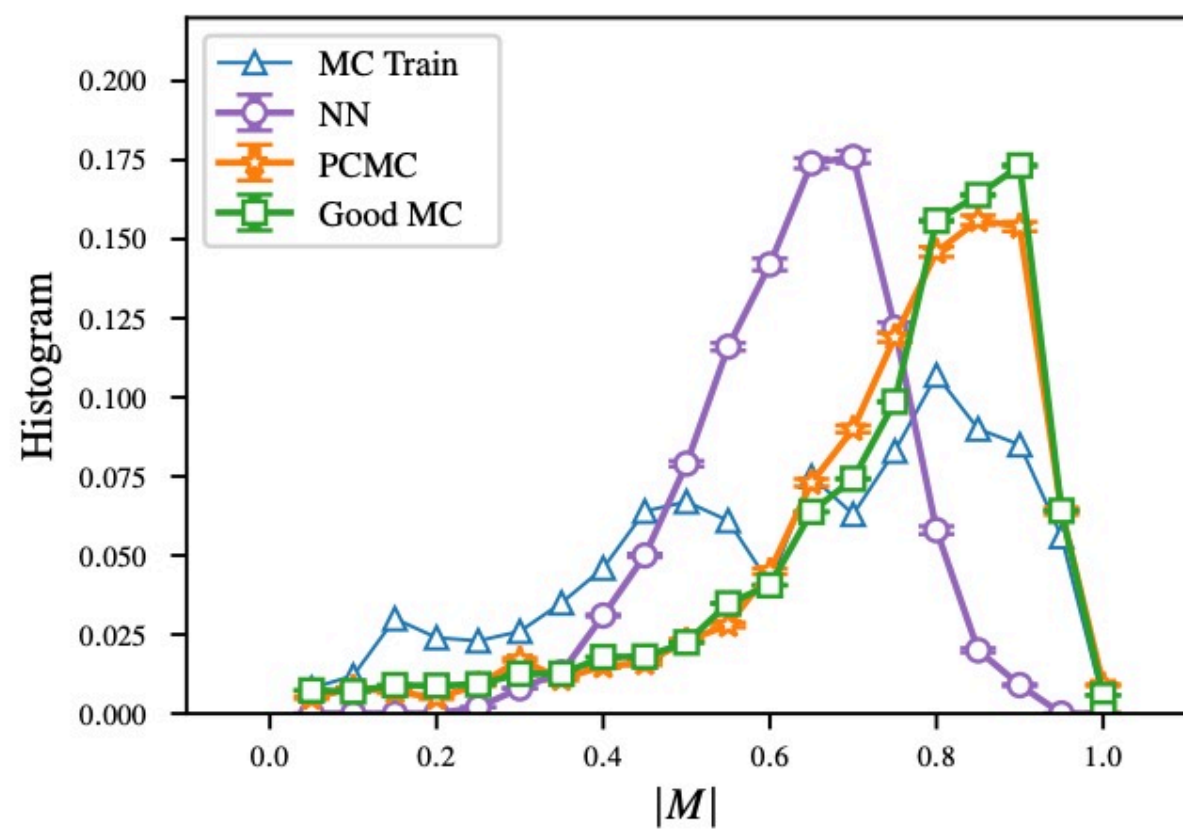
where \mathbf{G}_l is the l -th generated configuration, M_l and E_l refer to the magnetization $\frac{1}{N} |\sum_j \sigma_j|$ and the energy density $\frac{1}{N} \langle H \rangle$

(a) Square-lattice Ising Model at $T_c = 2.269 J$



(b) Honeycomb-lattice Hubbard Model at $U_c/t = 3.83$





Recurrent neural networks (RNNs)

Generative models

The ability to learn hierarchies of concepts, building up multiple layers of abstraction, seems to be fundamental to making sense of the world.

"basic research is what I'm doing when I don't know what I'm doing"

As the old joke goes, if you ask a scientist how far away some discovery is and they say "10 years" (or more), what they mean is "I've got no idea". AI, like controlled fusion and a few other technologies, has been 10 years away for 60 plus years.